

Jev vs ModernBERT-large

8-bucket answer-length prediction on WildChat

ModernBERT performs substantially better on every metric.

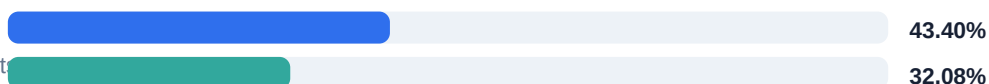
Comparable setup 45,001 test prompts • same bucket boundaries • 2,048-token input cap

■ ModernBERT-large

■ Jev 1.13

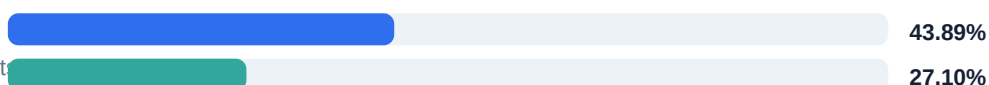
Exact-bucket accuracy

Jev difference: -11.32 percentage points



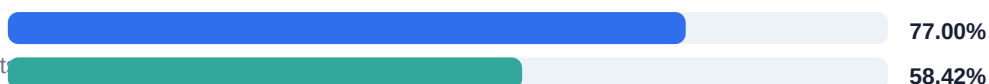
Macro-F1

Jev difference: -16.79 percentage points



Within one bucket

Jev difference: -18.58 percentage points



Conclusion

ModernBERT is the stronger predictor: +11.32 pp exact accuracy and +18.58 pp within-one-bucket accuracy. Jev still beats the prompt-length baseline (18.35% accuracy), but does not match the fine-tuned encoder.

Jev evaluation 42.17M input tokens • approximately \$1.77 • 0 API errors • 17.2 minutes